

Lab 2: Building a Real-Time User Activity Tracker with Kafka

Scenario

You are working for a fintech startup that needs to track user interactions in real-time. Your task is to implement a Kafka-based activity tracking system using the Kafka Java API.

Software Required

- Apache Kafka 3.0+ installed and running locally
 - JDK 11+
 - Maven or Gradle
 - Eclipse, IntelliJ IDEA or any Java IDE
-

Instructions

Step 1: Create a Kafka Topic

Open a Command Prompt and create a topic called user-interactions.

For Windows:

```
%KAFKA_HOME%\bin\windows\kafka-topics.bat --create --topic user-interactions  
^  
--bootstrap-server localhost:9092 --partitions 1 --replication-factor 1
```

Or use the provided script:

```
cd "Lab Code\Lab2\scripts"  
create-topic.bat
```

Step 2: Create the Producer Application

Create a Java class `UserActivityProducer.java` and implement a Kafka producer that sends the following types of events:

- "login"
- "view_profile"
- "transfer_initiated"

Each event should contain a `userId`, `action`, and `timestamp`.

Step 3: Create the Consumer Application

Create a Java class `UserActivityConsumer.java` that consumes from the `user-interactions` topic and logs each event to the console.

Step 4: Run the Producer and Consumer

Option A: Using Eclipse (Recommended) 1. Start the consumer in Eclipse: Right-click `UserActivityConsumer.java` → Run As → Java Application 2. Run the producer in Eclipse: Right-click `UserActivityProducer.java` → Run As → Java Application 3. Observe the messages being consumed in real-time.

Option B: Using Windows Scripts 1. Start the consumer: Open Command Prompt → Navigate to `Lab Code\Lab2\scripts` → Run `run-consumer.bat` 2. Run the producer: Open another Command Prompt → Navigate to `Lab Code\Lab2\scripts` → Run `run-producer.bat` 3. Observe the messages being consumed in real-time.

Option C: Manual Command Line 1. Start the consumer in one Command Prompt window. 2. Run the producer in another Command Prompt window. 3. Observe the messages being consumed in real-time.

Expected Output

When running the consumer, you should see output similar to:

```
Received message: key = 12345, value = {"userId": "12345", "action": "login",
"timestamp": "2025-07-18T10:00:00Z"}, partition = 0, offset = 0
Received message: key = 67890, value = {"userId": "67890", "action":
"transfer_initiated", "timestamp": "2025-07-18T10:01:00Z"}, partition = 0,
offset = 1
```

Complete Implementation

Project Structure

The complete project is available in the `Lab Code/Lab2/` folder with the following structure:

```
Lab2/
├── pom.xml                # Maven configuration
├── README.md              # Setup and usage guide
├── ECLIPSE_SETUP.md      # Eclipse-specific setup
├── src/
│   ├── main/java/com/fintech/
│   │   ├── UserActivityProducer.java    # Kafka Producer
│   │   ├── UserActivityConsumer.java    # Kafka Consumer
│   │   └── UserActivityEvent.java       # Event model class
├── scripts/
│   ├── create-topic.bat    # Windows script to create Kafka topic
│   ├── run-producer.bat   # Windows script to run producer
│   └── run-consumer.bat   # Windows script to run consumer
```

	start-kafka.bat	# Windows script to start Kafka services
	stop-kafka.bat	# Windows script to stop Kafka services

File: UserActivityProducer.java

```
import org.apache.kafka.clients.producer.*; import
org.apache.kafka.common.serialization.StringSerializer; import java.util.Properties;
import java.util.Random;

public class UserActivityProducer { public static void main(String[] args) throws
InterruptedException {

    Properties props = new Properties();

    props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
    props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());

        Producer<String, String> producer = new KafkaProducer<>(props);

        String[] actions = {"login", "view_profile", "transfer_initiated"};

        Random rand = new Random();

        for (int i = 0; i < 10; i++) {
            String userId = String.format("%05d", rand.nextInt(100000));
            String action = actions[rand.nextInt(actions.length)];
            String timestamp = java.time.LocalDateTime.now().toString();

            String json = String.format("{\"userId\": \"%s\", \"action\": \"%s\",
\\\"timestamp\\\": \"%s\"}",
                userId, action, timestamp);

            ProducerRecord<String, String> record = new ProducerRecord<>("user-
interactions", userId, json);
            producer.send(record, (metadata, exception) -> {
                if (exception != null) {
                    System.err.println("Error sending message: " +
exception.getMessage());
                } else {
                    System.out.printf("Sent: %s\\n", json);
                }
            });

            Thread.sleep(1000); // Simulate real-time delay
        }
    }
}
```

```
        producer.close();
    }
}
```

File: `UserActivityConsumer.java`

```
import org.apache.kafka.clients.consumer.*;
import org.apache.kafka.common.serialization.StringDeserializer;
import java.time.Duration;
import java.util.Collections;
import java.util.Properties;

public class UserActivityConsumer {
    public static void main(String[] args) {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "activity-consumer-group");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class.getName());
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class.getName());
        props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

        Consumer<String, String> consumer = new KafkaConsumer<>(props);
        consumer.subscribe(Collections.singletonList("user-interactions"));

        try {
            while (true) {
                ConsumerRecords<String, String> records =
consumer.poll(Duration.ofMillis(100));
                for (ConsumerRecord<String, String> record : records) {
                    System.out.printf("Received: %s\n", record.value());
                }
            }
        } finally {
            consumer.close();
        }
    }
}
```

Windows-Specific Setup

Prerequisites for Windows:

1. **Apache Kafka** installed and KAFKA_HOME environment variable set
2. **JDK 11+** installed and JAVA_HOME environment variable set
3. **Maven 3.6+** installed and added to PATH
4. **Eclipse IDE** with Maven plugin

Quick Start for Windows:

1. **Start Kafka Services:**

```
cd "Lab Code\Lab2\scripts"  
start-kafka.bat
```

2. **Create Topic:**

```
create-topic.bat
```

3. **Run Applications:**

- **Consumer:** run-consumer.bat
- **Producer:** run-producer.bat (in another Command Prompt)

Environment Variables:

Set the following environment variables in Windows: - KAFKA_HOME=C:\path\to\kafka - JAVA_HOME=C:\Program Files\Java\jdk-11 - Add %KAFKA_HOME%\bin\windows to PATH

Troubleshooting Windows Issues:

1. **Kafka not starting:** Check if KAFKA_HOME is set correctly
2. **Java not found:** Verify JAVA_HOME and PATH settings
3. **Permission denied:** Run Command Prompt as Administrator
4. **Port conflicts:** Ensure ports 2181 (Zookeeper) and 9092 (Kafka) are available